

```
In [18]: import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt
```

```
In [19]: IMAGE_SIZE=256
BATCH_SIZE=32
CHANNELS=3
EPCOHS=30
```

```
In [20]: dataset=tf.keras.preprocessing.image_dataset_from_directory(
    "PlantVillage",
    shuffle=True,
    image_size = (IMAGE_SIZE,IMAGE_SIZE),
    batch_size = BATCH_SIZE
)
```

Found 2152 files belonging to 3 classes.

```
In [21]: class_names=dataset.class_names
class_names
```

```
Out[21]: ['Potato __ Early_blight', 'Potato __ Late_blight', 'Potato __ healthy']
```

```
In [22]: len(dataset) #shows 68 bcz every element in dataset is a batch of 32 images
```

```
Out[22]: 68
```

```
In [23]: 68*32 #Last batch not perfect so shows a bit more than that, 68 BATCHES OF 32 IMAGE
```

```
Out[23]: 2176
```

```
In [24]: for image_batch, labels_batch in dataset.take(1):
    print(image_batch.shape)
    print(labels_batch.numpy())
(32, 256, 256, 3)
[2 1 1 0 1 0 0 0 1 2 2 0 2 1 0 1 1 1 0 0 1 2 0 0 0 2 0 1 2 1 1 1]
```

```
In [80]: plt.figure(figsize=(10, 10))
for image_batch, labels_batch in dataset.take(1):
    for i in range(12):
        ax = plt.subplot(3, 4, i + 1)
        plt.imshow(image_batch[i].numpy().astype("uint8"))
        plt.title(class_names[labels_batch[i]])
        plt.axis("off")
```

Potato_Late_blight



Potato_healthy



Potato_Late_blight



Potato_Late_blight



Potato_Late_blight



Potato_Late_blight



Potato_Late_blight



Potato_Early_blight



Potato_Late_blight



Potato_Early_blight



Potato_Late_blight



Potato_Early_blight



```
In [ ]: Function to Split Dataset
Dataset should be bifurcated into 3 subsets, namely:
```

1. Training: Dataset to be used while training 80%
2. Validation: Dataset to be tested against while training 10%
3. Test: Dataset to be tested against after we trained a model 10%

```
In [27]: train_size=0.8
len(dataset)*train_size
```

Out[27]: 54.400000000000006

```
In [28]: train_ds=dataset.take(54)
len(train_ds)
```

Out[28]: 54

```
In [29]: test_ds=dataset.skip(54)
len(test_ds)
```

Out[29]: 14

```
In [30]: val_size=0.1
len(dataset)*val_size
```

Out[30]: 6.800000000000001

```
In [31]: val_ds=test_ds.take(6)
len(val_ds)
```

Out[31]: 6

```
In [32]: test_ds = test_ds.skip(6)
len(test_ds)
```

Out[32]: 8

```
In [33]: def get_dataset_partitions_tf(ds, train_split=0.8, val_split=0.1, test_split=0.1, shuffle=False):
    assert (train_split + test_split + val_split) == 1

    ds_size = len(ds)

    if shuffle:
        ds = ds.shuffle(shuffle_size, seed=12)

    train_size = int(train_split * ds_size)
    val_size = int(val_split * ds_size)

    train_ds = ds.take(train_size)
    val_ds = ds.skip(train_size).take(val_size)
    test_ds = ds.skip(train_size).skip(val_size)

    return train_ds, val_ds, test_ds
```

```
In [34]: train_ds, val_ds, test_ds = get_dataset_partitions_tf(dataset)
```

```
In [35]: len(train_ds)
```

Out[35]: 54

```
In [36]: len(val_ds)
```

Out[36]: 6

```
In [37]: len(test_ds)
```

Out[37]: 8

```
In [ ]: Cache, Shuffle, and Prefetch the Dataset #Improves performance by reading and keeping data in memory
```

```
In [38]: train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
```

```
test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
```

In [39]:

```
resize_and_rescale = tf.keras.Sequential([
    layers.Resizing(IMAGE_SIZE, IMAGE_SIZE),
    layers.Rescaling(1.0/255),
])
```

In []: Data Augmentation

In [45]:

```
data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal_and_vertical"),
    layers.RandomRotation(0.2),
])
```

In []: Model Architecture || Building model
 We use a CNN coupled **with** a Softmax activation **in** the output layer.
 We also add the initial layers **for** resizing, normalization **and** Data Augmentation.

In [47]:

```
model = models.Sequential([
    layers.Input(shape=(IMAGE_SIZE, IMAGE_SIZE, CHANNELS)),
    resize_and_rescale,

    layers.Conv2D(32, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(3, activation='softmax')
])
```

In [48]:

```
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape
sequential (Sequential)	(None, 256, 256, 3)
conv2d_12 (Conv2D)	(None, 254, 254, 32)
max_pooling2d_12 (MaxPooling2D)	(None, 127, 127, 32)
conv2d_13 (Conv2D)	(None, 125, 125, 64)
max_pooling2d_13 (MaxPooling2D)	(None, 62, 62, 64)
conv2d_14 (Conv2D)	(None, 60, 60, 64)
max_pooling2d_14 (MaxPooling2D)	(None, 30, 30, 64)
conv2d_15 (Conv2D)	(None, 28, 28, 64)
max_pooling2d_15 (MaxPooling2D)	(None, 14, 14, 64)
flatten_3 (Flatten)	(None, 12544)
dense_3 (Dense)	(None, 64)
dense_4 (Dense)	(None, 3)

Total params: 896,323 (3.42 MB)

Trainable params: 896,323 (3.42 MB)

Non-trainable params: 0 (0.00 B)

In []: Compiling the Model
We use adam Optimizer, SparseCategoricalCrossentropy for losses, accuracy as a metric

In [49]: model.compile(
optimizer='adam',
loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
metrics=['accuracy'])

In [50]: history = model.fit(
train_ds,
batch_size=BATCH_SIZE,
validation_data=val_ds,
verbose=1,
epochs=30,
)

```
Epoch 1/30
54/54 122s 2s/step - accuracy: 0.5041 - loss: 0.9183 - val_accuracy: 0.8177 - val_loss: 0.6340
Epoch 2/30
54/54 132s 2s/step - accuracy: 0.8608 - loss: 0.4147 - val_accuracy: 0.8698 - val_loss: 0.2926
Epoch 3/30
54/54 114s 2s/step - accuracy: 0.8939 - loss: 0.2913 - val_accuracy: 0.9115 - val_loss: 0.1765
Epoch 4/30
54/54 92s 2s/step - accuracy: 0.9461 - loss: 0.1422 - val_accuracy: 0.9583 - val_loss: 0.1477
Epoch 5/30
54/54 87s 2s/step - accuracy: 0.9423 - loss: 0.1592 - val_accuracy: 0.8854 - val_loss: 0.3076
Epoch 6/30
54/54 95s 2s/step - accuracy: 0.9426 - loss: 0.1392 - val_accuracy: 0.9583 - val_loss: 0.1021
Epoch 7/30
54/54 92s 2s/step - accuracy: 0.9666 - loss: 0.0866 - val_accuracy: 0.9844 - val_loss: 0.0404
Epoch 8/30
54/54 90s 2s/step - accuracy: 0.9625 - loss: 0.0945 - val_accuracy: 0.9844 - val_loss: 0.0457
Epoch 9/30
54/54 131s 2s/step - accuracy: 0.9805 - loss: 0.0589 - val_accuracy: 0.9844 - val_loss: 0.0334
Epoch 10/30
54/54 125s 2s/step - accuracy: 0.9883 - loss: 0.0336 - val_accuracy: 0.9531 - val_loss: 0.1552
Epoch 11/30
54/54 111s 2s/step - accuracy: 0.9652 - loss: 0.0899 - val_accuracy: 0.9792 - val_loss: 0.0489
Epoch 12/30
54/54 94s 2s/step - accuracy: 0.9872 - loss: 0.0439 - val_accuracy: 0.9896 - val_loss: 0.0248
Epoch 13/30
54/54 98s 2s/step - accuracy: 0.9980 - loss: 0.0070 - val_accuracy: 1.0000 - val_loss: 0.0073
Epoch 14/30
54/54 156s 3s/step - accuracy: 0.9996 - loss: 0.0045 - val_accuracy: 0.9844 - val_loss: 0.0323
Epoch 15/30
54/54 105s 2s/step - accuracy: 0.9977 - loss: 0.0082 - val_accuracy: 0.9948 - val_loss: 0.0148
Epoch 16/30
54/54 99s 2s/step - accuracy: 0.9974 - loss: 0.0109 - val_accuracy: 0.9948 - val_loss: 0.0087
Epoch 17/30
54/54 113s 2s/step - accuracy: 1.0000 - loss: 0.0011 - val_accuracy: 0.9948 - val_loss: 0.0082
Epoch 18/30
54/54 142s 3s/step - accuracy: 1.0000 - loss: 3.4857e-04 - val_accuracy: 0.9948 - val_loss: 0.0049
Epoch 19/30
54/54 84s 2s/step - accuracy: 1.0000 - loss: 2.3783e-04 - val_accuracy: 0.9948 - val_loss: 0.0049
```

```

accuracy: 0.9948 - val_loss: 0.0049
Epoch 20/30
54/54 ━━━━━━━━━━ 79s 1s/step - accuracy: 1.0000 - loss: 1.9179e-04 - val_a
ccuracy: 0.9948 - val_loss: 0.0050
Epoch 21/30
54/54 ━━━━━━━━━━ 78s 1s/step - accuracy: 1.0000 - loss: 1.7389e-04 - val_a
ccuracy: 0.9948 - val_loss: 0.0054
Epoch 22/30
54/54 ━━━━━━━━━━ 81s 1s/step - accuracy: 1.0000 - loss: 1.5856e-04 - val_a
ccuracy: 0.9948 - val_loss: 0.0065
Epoch 23/30
54/54 ━━━━━━━━━━ 80s 1s/step - accuracy: 1.0000 - loss: 1.0653e-04 - val_a
ccuracy: 0.9948 - val_loss: 0.0061
Epoch 24/30
54/54 ━━━━━━━━━━ 78s 1s/step - accuracy: 1.0000 - loss: 1.3264e-04 - val_a
ccuracy: 0.9948 - val_loss: 0.0063
Epoch 25/30
54/54 ━━━━━━━━━━ 77s 1s/step - accuracy: 1.0000 - loss: 1.0586e-04 - val_a
ccuracy: 0.9948 - val_loss: 0.0056
Epoch 26/30
54/54 ━━━━━━━━━━ 79s 1s/step - accuracy: 1.0000 - loss: 8.7578e-05 - val_a
ccuracy: 0.9948 - val_loss: 0.0069
Epoch 27/30
54/54 ━━━━━━━━━━ 78s 1s/step - accuracy: 1.0000 - loss: 9.0309e-05 - val_a
ccuracy: 0.9948 - val_loss: 0.0056
Epoch 28/30
54/54 ━━━━━━━━━━ 78s 1s/step - accuracy: 1.0000 - loss: 5.8491e-05 - val_a
ccuracy: 0.9948 - val_loss: 0.0078
Epoch 29/30
54/54 ━━━━━━━━━━ 77s 1s/step - accuracy: 1.0000 - loss: 6.2423e-05 - val_a
ccuracy: 0.9948 - val_loss: 0.0067
Epoch 30/30
54/54 ━━━━━━━━━━ 79s 1s/step - accuracy: 1.0000 - loss: 5.1099e-05 - val_a
ccuracy: 0.9948 - val_loss: 0.0078

```

In [51]: `scores = model.evaluate(test_ds)`

```
8/8 ━━━━━━━━━━ 8s 318ms/step - accuracy: 0.9953 - loss: 0.0063
```

In []: Achieved 99.53% accuracy on our test dataset

In [52]: `scores`

Out[52]: [0.0052637276239693165, 0.99609375]

In []: Plotting the Accuracy and Loss Curves

In [53]: `history.params`

Out[53]: {'verbose': 1, 'epochs': 30, 'steps': 54}

In [54]: `history.history.keys()`

Out[54]: dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])

```
In [ ]: loss, accuracy, val loss etc are a python list containing values of loss, accuracy
```

```
In [55]: type(history.history['loss'])
```

```
Out[55]: list
```

```
In [56]: len(history.history['loss'])
```

```
Out[56]: 30
```

```
In [57]: history.history['loss'][:5] # show Loss for first 5 epochs
```

```
Out[57]: [0.7739138603210449,  
          0.3689870536327362,  
          0.2585057318210602,  
          0.16928669810295105,  
          0.13185274600982666]
```

```
In [58]: acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']  
  
loss = history.history['loss']  
val_loss = history.history['val_loss']
```

```
In [61]: plt.figure(figsize=(8, 8))  
plt.subplot(1, 2, 1)  
plt.plot(range(30), acc, label='Training Accuracy')  
plt.plot(range(30), val_acc, label='Validation Accuracy')  
plt.legend(loc='lower right')  
plt.title('Training and Validation Accuracy')  
  
plt.subplot(1, 2, 2)  
plt.plot(range(30), loss, label='Training Loss')  
plt.plot(range(30), val_loss, label='Validation Loss')  
plt.legend(loc='upper right')  
plt.title('Training and Validation Loss')  
plt.show()
```



RUNNING PREDICTION ON A SAMPLE IMAGE

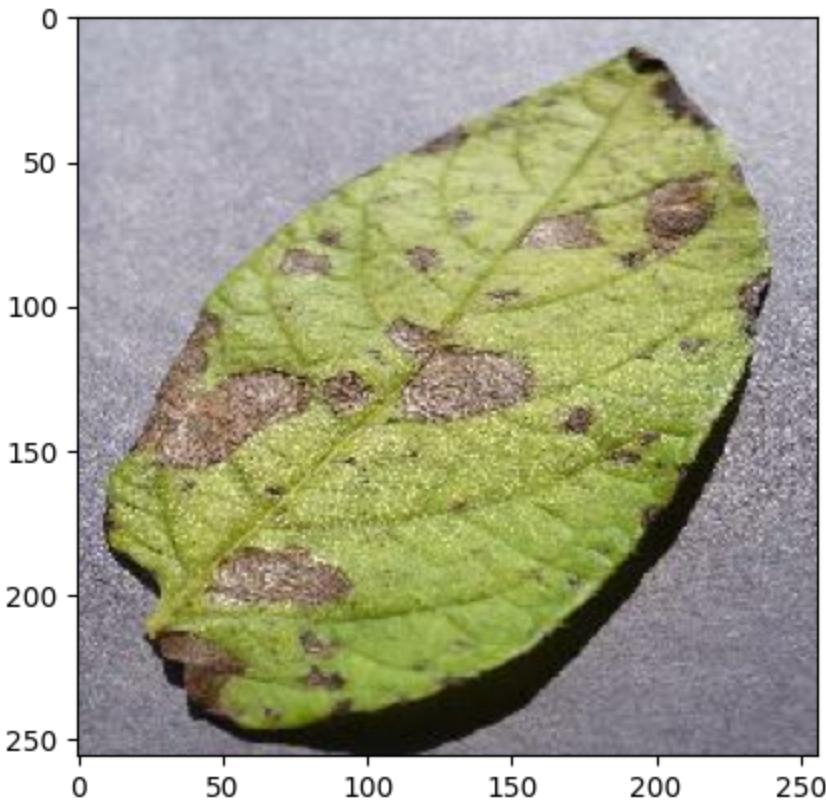
```
In [93]: import numpy as np
for images_batch, labels_batch in test_ds.take(1):

    first_image = images_batch[0].numpy().astype('uint8')
    first_label = labels_batch[0].numpy()

    print("first image to predict")
    plt.imshow(first_image)
    print("actual label:", class_names[first_label])

    batch_prediction = model.predict(images_batch)
    print("predicted label:", class_names[np.argmax(batch_prediction[0])])

first image to predict
actual label: Potato__Early_blight
1/1 ━━━━━━ 0s 448ms/step
predicted label: Potato__Early_blight
```



In []: FUNCTION FOR INFERENCE

```
In [63]: def predict(model, img):
    img_array = tf.keras.preprocessing.image.img_to_array(images[i].numpy())
    img_array = tf.expand_dims(img_array, 0)

    predictions = model.predict(img_array)

    predicted_class = class_names[np.argmax(predictions[0])]
    confidence = round(100 * (np.max(predictions[0])), 2)
    return predicted_class, confidence
```

In []: RUNNING INFERENCE ON A FEW SAMPLE IMAGES

```
In [92]: plt.figure(figsize=(15, 15))
for images, labels in test_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))

        predicted_class, confidence = predict(model, images[i].numpy())
        actual_class = class_names[labels[i]]

        plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class}.\n Confi
        plt.axis("off")
```

```
1/1 _____ 0s 127ms/step
1/1 _____ 0s 94ms/step
1/1 _____ 0s 91ms/step
1/1 _____ 0s 107ms/step
1/1 _____ 0s 202ms/step
1/1 _____ 0s 125ms/step
1/1 _____ 0s 140ms/step
1/1 _____ 0s 152ms/step
1/1 _____ 0s 151ms/step
```

Actual: Potato__Late_blight,
Predicted: Potato__Late_blight.
Confidence: 100.0%



Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 100.0%

Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 100.0%



Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 100.0%



Actual: Potato__healthy,
Predicted: Potato__healthy.
Confidence: 100.0%



Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 100.0%



Actual: Potato__Late_blight,
Predicted: Potato__Late_blight.
Confidence: 100.0%



```
In [66]: model.save('my_model.keras')
```

```
In [83]: print(model)
<Sequential name=sequential_3, built=True>
```

```
In [79]: import os
```

```
models_dir = ".../saved_model"

if not os.path.exists(models_dir):
    os.makedirs(models_dir)

model_version = max([int(i.split('.')[0]) for i in os.listdir(models_dir) if i.endswith('.h5')])

model.save(f"{models_dir}/{model}_{model_version}.keras")
```

In [85]:

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report, recall_score
```

In [86]:

```
y_true = []
y_pred = []

for images, labels in test_ds:
    preds = model.predict(images)
    preds = np.argmax(preds, axis=1)
    y_true.extend(labels.numpy())
    y_pred.extend(preds)

y_true = np.array(y_true)
y_pred = np.array(y_pred)
```

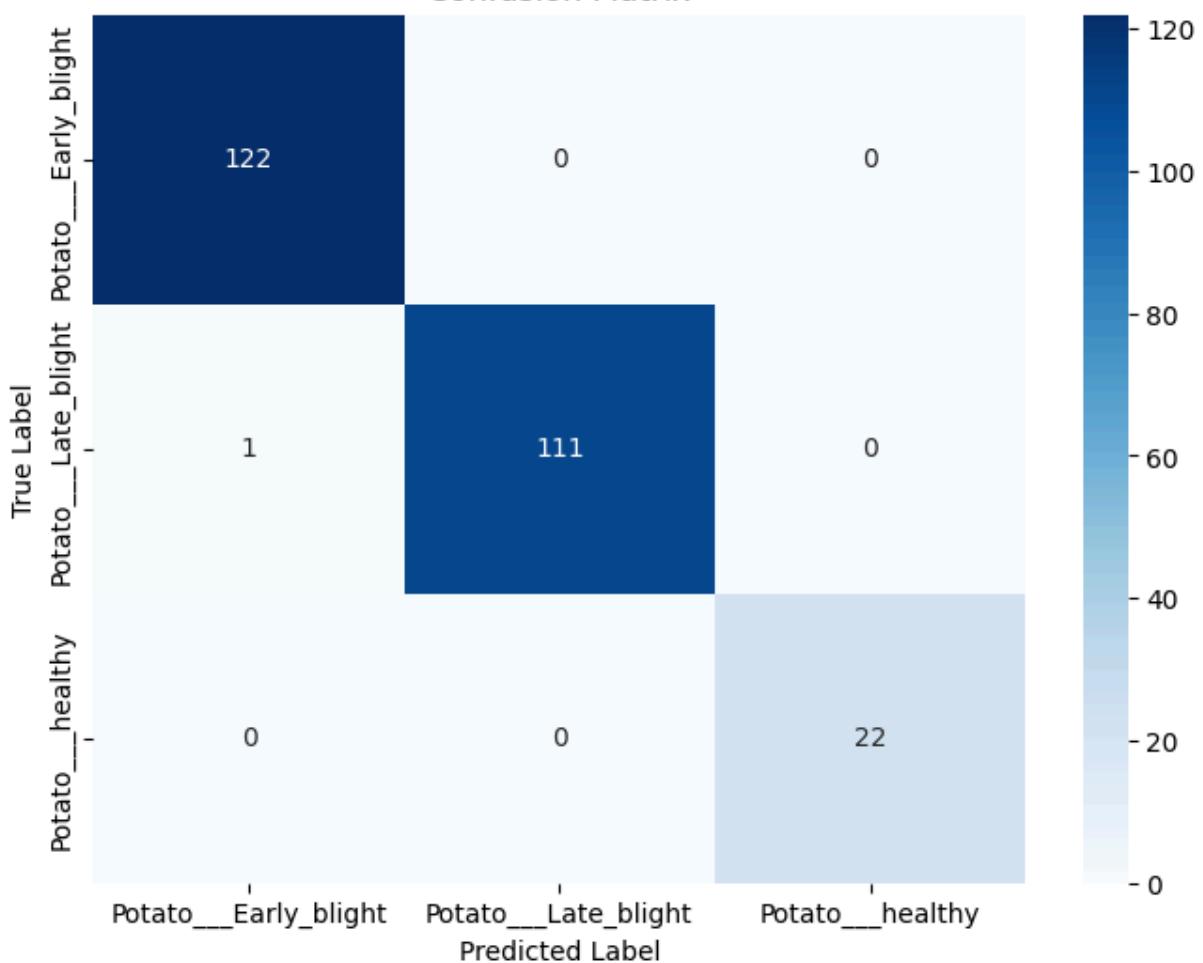
```
1/1 ━━━━━━━━ 2s 2s/step
1/1 ━━━━━━ 1s 671ms/step
1/1 ━━━━ 1s 616ms/step
1/1 ━━ 1s 564ms/step
1/1 ━ 1s 555ms/step
1/1 0s 494ms/step
1/1 ━ 1s 584ms/step
1/1 ━━ 1s 521ms/step
```

In [87]:

```
cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

Confusion Matrix



```
In [88]: print(classification_report(y_true, y_pred, target_names=class_names))
```

	precision	recall	f1-score	support
Potato_Early_blight	0.99	1.00	1.00	122
Potato_Late_blight	1.00	0.99	1.00	112
Potato_healthy	1.00	1.00	1.00	22
accuracy			1.00	256
macro avg	1.00	1.00	1.00	256
weighted avg	1.00	1.00	1.00	256

```
In [89]: recall_macro = recall_score(y_true, y_pred, average='macro')
recall_per_class = recall_score(y_true, y_pred, average=None)

print(f"Macro Average Recall: {recall_macro:.4f}\n")
for idx, class_name in enumerate(class_names):
    print(f"Recall for {class_name}: {recall_per_class[idx]:.4f}")
```

Macro Average Recall: 0.9970

Recall for Potato_Early_blight: 1.0000
 Recall for Potato_Late_blight: 0.9911
 Recall for Potato_healthy: 1.0000

```
In [90]: model.save('my_model.keras')
```

```
In [ ]:
```